



COMP9517

Computer Vision

Introduction

What is Computer Vision?

Every picture tells a story



Computer vision automates and integrates many information processing and representation approaches useful for visual perception

What is computer vision?

Computer science perspective

Computer vision is the **interdisciplinary field** that develops **theories and methods** to allow computers **extract relevant information from digital images or videos**

Computer engineering perspective

Computer vision is the **interdisciplinary field** that develops **algorithms and tools** to **automate perceptual tasks** normally performed by the human visual system


Can computers match (or beat) human vision?



Yes and no (but mostly no)



- Humans are much better at “hard” tasks
- Computers can be better at “easy” tasks

Human vision has its limitations...




Which objects are brighter?

Human vision has its limitations...





Which objects are brighter?

Human vision has its limitations...





Which side of this object is brighter?

Human vision has its limitations...




Which side of this object is brighter?

Human vision has its limitations...




Which side of this object is brighter?

Human vision has its limitations...




Are the cells popping in or out?

Human vision has its limitations...




Are the cells popping in or out?

Human vision has its limitations...



What is the pattern
you see?

Human vision has its limitations...



What is the pattern
you see?

Human vision has its limitations...




What object do you
see in this image?

Human vision has its limitations...




What object do you
see in this image?

Human vision has its limitations...




Are the main lines
parallel??

Human vision has its limitations...



Are the main lines
parallel?

Human vision has its limitations...



In which direction are these particles moving ?

Human vision has its limitations...




<https://www.youtube.com/watch?v=a7efEggplrE>

Course rationale

Human vision has its limitations

- intensities, shapes, patterns, motions can be misinterpreted
- it is labor intensive, time-consuming, subjective, error-prone



Computer vision can potentially improve this

- work day and night without getting tired
- analyse information quantitatively and objectively
- potentially more accurate, precise, reproducible

If the methods and tools are well designed!



Application: 3D shape reconstruction

Project [VarCity](#) recreates 3D city models using social media photos




Application: image classification and captioning

[Google's Show and Tell](#) open-source image captioning model in TensorFlow

Human captions from the training set



Automatically captioned




Application: intelligent collision avoidance

[Iris Automation](#) provides safer drone operation with intelligent collision avoidance



Application: face detection and recognition

Facebook's [DeepFace](#) project nears human accuracy in identifying faces




Application: face detection and recognition

For improving image capture on digital cameras





Application: vision-based biometrics



[How the Afghan girl was identified by her iris patterns](#)

The remarkable story of Sharbat Gula, first photographed in 1984 aged 12 in a refugee camp in Pakistan by National Geographic photographer Steve McCurry, and traced 18 years later to a remote part of Afghanistan where she was again photographed by McCurry...



Application: logging in without a password



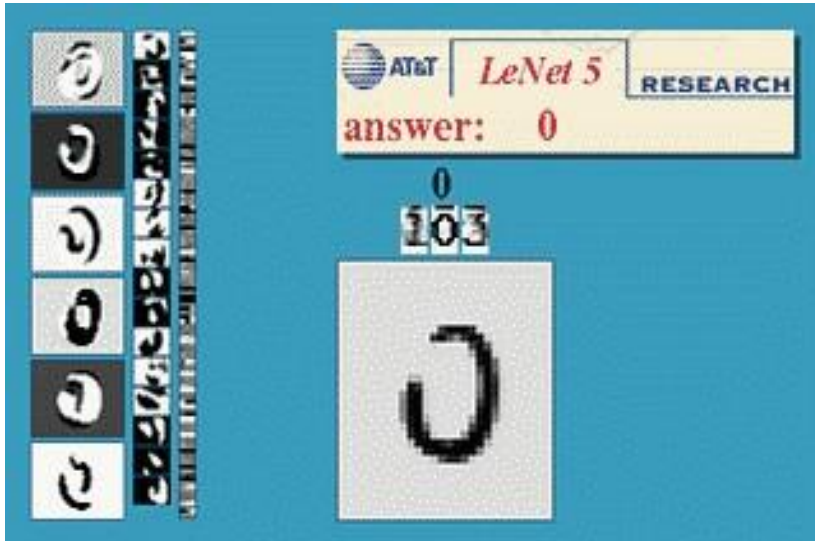
Fingerprint scanners on
modern laptops and
other devices



Windows Hello makes
logging in as easy as
looking at your PC

Application: optical character recognition (OCR)

Converting scanned documents or number plates to processable text



Application: object recognition in supermarkets



[LaneHawk by Evolution Robotics Retail](#)
provides a loss-prevention solution that
detects bottom-of-basket (BOB) items
in checkout lanes

Application: object recognition in phones



Application: autonomous vehicles

[Intel's Mobileye](#) makes cars safer and more autonomous

The screenshot displays the Intel Mobileye website with a focus on autonomous vehicle technology. The main banner features a top-down view of a car with three yellow beams representing camera fields of view, labeled 'rear looking camera', 'forward looking camera', and 'side looking camera'. The headline reads 'Our Vision. Your Safety.' with navigation tabs for 'manufacturer products' and 'consumer products'.

Below the banner are three product highlights:

- EyeQ Vision on a Chip:** Accompanied by an image of the EyeQ chip and a 'read more' link.
- Vision Applications:** Includes an image of a pedestrian on a crosswalk and the text 'Road, Vehicle, Pedestrian Protection and more', with a 'read more' link.
- AWS Advance Warning System:** Features an image of a circular display showing a car icon and the number '0.8', with a 'read more' link.

On the right side, there are two vertical sections:

- News:** Contains two headlines: 'Mobileye Advanced Technologies Power Volvo Cars World First Collision Warning With Auto Brake System' and 'Volvo: New Collision Warning with Auto Brake Helps Prevent Rear-end', followed by an 'all news' link.
- Events:** Lists 'Mobileye at Equip Auto, Paris, France' and 'Mobileye at SEMA, Las Vegas, NV', with a 'read more' link at the bottom.

Application: space exploration


NASA's Mars Exploration Rover Spirit autonomously captured this picture in 2007




Vision systems used for panorama stitching, 3D terrain modeling, obstacle detection, position tracking

See [Computer Vision on Mars](#) for more information

Application: machine vision in robotics




[NASA's Mars Spirit Rover](#)



[RoboCup](#)

Application: medical imaging




Computer Aided Diagnosis



Image Guided Surgery

Application: video surveillance



- Traffic monitoring
- Person tracking
- Action recognition
- Speed estimation
- Object counting
- ...



Goals of Computer Vision

- Extract useful information from images: both **metric** and **semantic**
- **Complexity** of visual data is a challenge
- Recent progress due to higher processing power, memory, storage capacity
- Image→measurements→model→algorithms for learning and inference

Computer vision tasks

- Obtain simple inferences from individual pixel values
- Group pixels to separate object regions or infer shape information
- Recognise objects using geometric or statistical pixel information
- Combine information from multiple images into a coherent whole

Requires understanding of the **physics of imaging** and the use of **mathematical and statistical models** for information extraction

Critical issues in computer vision

- **Sensing:** how do sensors obtain images of the world?
- **Encoded Information:** how do images yield information of the scene, such as colour, texture, shape, motion...?
- **Representations:** what representations are appropriate to describe objects?
- **Algorithms:** what are appropriate algorithms to process image information and construct scene descriptions?

Low-level computer vision

This is almost entirely **digital image processing** (image in > image out)

- **Sensing:** image capture and digitisation
- **Preprocessing:** suppress noise and enhance object features
- **Segmentation:** separate objects from background and partition them
- **Description:** compute features that differentiate objects
- **Classification:** assign labels to image segments (regions)

High-level computer vision

This is about knowledge construction, representation, and inference

- **Recognition:** identify objects based on low-level information
- **Interpretation:** assign meaning to groups of recognized objects
- **Scene analysis:** complete understanding of the captured scene

Assumed knowledge

To do this course successfully you should:

- Be able to program well in **Python** or willing to learn it independently
- Be familiar with **data structures and algorithms** and basic **statistics**
- Be able/learn to **use and integrate software packages** (OpenCV, Scikit-Learn, Keras)
- Be familiar with **vector calculus and linear algebra** or willing to learn it independently

*Please self-assess **before** deciding to stay/enroll in the course*

Student learning outcomes

After completing this course you will be able to:

- Explain basic scientific, statistical, and engineering approaches to computer vision
- Implement and test computer vision algorithms using existing software platforms
- Build larger computer vision applications by integrating software modules
- Interpret and comment on articles in the computer vision literature

Course Changes in T3, 2020

Based on feedback received in T1 and T2, 2020:

1. All course components are **entirely online**. **Live online lectures** will provide opportunity to interact with lecturer. Online labs will be interactive.
2. A **new lab** to help prepare for online exam, **will not be assessed.**; also reduces proportion of marks for writing style assessments.
3. All labs and lecture slides will be **released as early as possible**.
4. The **level of work for group project** commensurate to resources available for an online course.
5. The group project will also have an **individual component**, as in T1, 2020.
6. The **exam** will consist of **critiquing a paper** from literature based on what has been learned in the course.

Weekly Class Structure

Week	Topic	Lecturer
1	Introduction & Image Formation	Professor Arcot Sowmya, Dr Wafa Johal
2	Image Processing	Dr Wafa Johal
3	Feature Representation	Dr Wafa Johal
4	Image Segmentation	Professor Arcot Sowmya
5	Image Segmentation, Motion Tracking	Professor Arcot Sowmya
6	Flexible Week (No Lectures, consultations will be held)	
7	Pattern Recognition	Dr Gelareh Mohammadi
8	Pattern Recognition, Deep Learning	Dr Gelareh Mohammadi
9	Applications	Professor Arcot Sowmya
10	Project Demos	Professor Arcot Sowmya, Dr Wafa Johal, Dr Gelareh Mohammadi

Weekly Class Structure

Week	Lecture Time Lecturer	Topic
Week 1	Monday (12-2 pm) Prof Arcot Sowmya Friday (12-2 pm) Dr Wafa Johal	Introduction Image Formation, Image Processing
Week 2	Monday (12-2 pm) Dr Wafa Johal Friday (12-1 pm) Dr Wafa Johal	Image processing
Week 3	Monday (12-2 pm) Dr Wafa Johal Friday (12-1 pm) Dr Wafa Johal	Feature representation
Week 4	Monday PUBLIC HOLIDAY Friday (12-2 pm) Prof Arcot Sowmya	Segmentation
Week 5	Monday (12-2 pm) Prof Arcot Sowmya Friday (12-1 pm) Prof Arcot Sowmya	Segmentation, Motion Tracking
Week 6	FLEXIBLE WEEK	No lectures; consultations will be held
Week 7	Monday (12-2 pm) Dr Gelareh Mohammadi Friday (12-1 pm) Dr Gelareh Mohammadi	Pattern Recognition
Week 8	Monday (12-2 pm) Dr Gelareh Mohammadi Friday (12-1 pm) Dr Gelareh Mohammadi	Pattern Recognition, Deep Learning
Week 9	Monday (12-2 pm) Prof Arcot Sowmya Friday (12-1 pm) Prof Arcot Sowmya	Applications
Week 10	Monday (1-3 pm) Friday (12-2 pm)	Project Demos

Weekly Class Structure

- **Lectures:** Mon and Fri, see previous slide for details; note PUBLIC HOLIDAY on Monday Oct 5th
- **Labs:** 1-2 PM on Fridays in weeks 2, 3, 5, 7, 8
- **Project consultations:**
 - On Fridays in weeks 6, 9
 - additional consultations require appointments with your assigned tutor
- **Project demo:**
 - On Monday and Friday in week 10, during class hours
 - Detailed schedule will be announced on class web page

ALL changes will be announced on class web page on WebCMS3

Assessments

Late Submission Penalty

Unless you have received special dispensation from the Lecturer in Charge, work that is submitted after the deadline DURING THE TERM will incur a penalty of 10% per day, up to a maximum of 100%.

For the final examination, university exam rules will apply.

Assessment	Marks	Release	Due
Assignment	10%	Week 1	Week 3
Lab Work	10%	Weeks 2, 3, 5, 7, 8	Weeks 3, 4, 6, 8, 9
Project (multiple stages)			
- Individual component	15%	Week 5	Week 7
- Group component *	25%	Week 5	Week 10
Exam	40%	Exam Period	Exam Period

Communication- Modes and Etiquette

- **Online forum (Piazza) is your first port of call-** post query of wider interest on lectures, labs, assessments
- **Contact LIC** for late submission, absence, assessment deadlines, lab and assessment content
- **Contact Course admin** for issues with enrolment, file submission, group enrolment or any admin matter
- every effort will be made to respond quickly to queries- allow **maximum of 24 hours turnaround**
- Do observe **standards of equity and respect** in dealing with all students and staff- in person, emails, forum posts, all other communication
- Preferred language of **communication** is **English**

Special Consideration/Supp Policy

- If your work in this course is affected by unforeseen adverse circumstances, you should apply for **Special Consideration**
- UNSW handles special consideration requests centrally. **Do not just email the LIC** about special consideration.
- Special Consideration **requests must be accompanied by documentation**
- Mark calculated in the same way as other students who sat the original assessment
- If you are awarded a Supp and do not attend, then your exam mark will be zero.

More information on Course web page

Plagiarism Policy

READ the UNSW Policy and Procedure on this.

For the **purposes of COMP 9517**, plagiarism includes copying or obtaining all, or a substantial part, of the material for your assignment, whether written or graphical report material, or software code, **without written acknowledgement** in your assignment from:

- a location on the Internet
- a book, article or other written document (published or unpublished) whether electronic or on paper or other medium
- another student, whether in your class or another class
- someone else (e.g. from someone who writes assignments for money)

Plagiarism (ctd)

- If you copy material from another student or non-student **with acknowledgement**, you will not be penalised for plagiarism, but the marks you get for this will be at the marker's discretion, and will reflect the marker's perception of the amount of work you put into finding and/or adapting the code/text.
- If you use text found in a publication (on the Internet or otherwise) then the marks you get for this will be at the **marker's discretion**, and will reflect the marker's **perception of the amount of work** you put into finding and/or adapting the text.

The assessments provide opportunities for you to develop important skills. Use these opportunities!

Further information on [WebCMS](#)

Please be sure you are familiar with:

- [Communication Etiquette](#)
- [Special Consideration](#)
- [Student Conduct](#)
- [Plagiarism Policy](#)
- [Academic Integrity](#)

Further reading on lecture topics

In the lectures we will be referring to various online resources for further reading such as:

- Richard Szeliski, [Computer Vision: Algorithms and Applications](#), Springer, 2011
- Dana H. Ballard and Christopher M. Brown, [Computer Vision](#), Prentice Hall, 1982
- Ian Goodfellow, Yoshua Bengio, Aaron Courville, [Deep Learning](#), MIT Press, 2016
- David A. Forsyth and Jean Ponce, [Computer Vision: A Modern Approach](#), Prentice Hall, 2011
- Simon J. D. Prince, [Computer Vision: Models, Learning and Inference](#), Cambridge University Press, 2012

And other books, scientific articles, and other resources available online or via the UNSW Library

Further reading on today's topics

- Chapter 1 of Szeliski for a general introduction to computer vision
- Chapter 1, Shapiro and Stockman
- Appendix A and B of Szeliski for background on linear algebra, numerical techniques and statistics

Acknowledgements

- Some images on applications taken from Szeliski with original sources credited where possible
- Other images and videos credited where possible