


# COMP9517: Computer Vision

Feature Representation

Part 1

# What are Features?



# Image Features

- Image features are essentially vectors that are a compact representation of images
- They represent important information shown in an image
- Intuitive examples of image features:
  - Blobs
  - Edges
  - Corners
  - Ridges
  - Circles
  - Ellipses
  - Lines
  - Etc...



# Image Features


- We need to represent images as feature vectors for further processing in a more efficient and robust way
- Examples of further processing include:
  - Object detection
  - Image segmentation
  - Image classification
  - Content-based image retrieval
  - Image stitching
  - Object tracking




# Object Detection



# Segmentation




# Image Classification






# Content-Based Image Retrieval







# Image Stitching



# Object Tracking



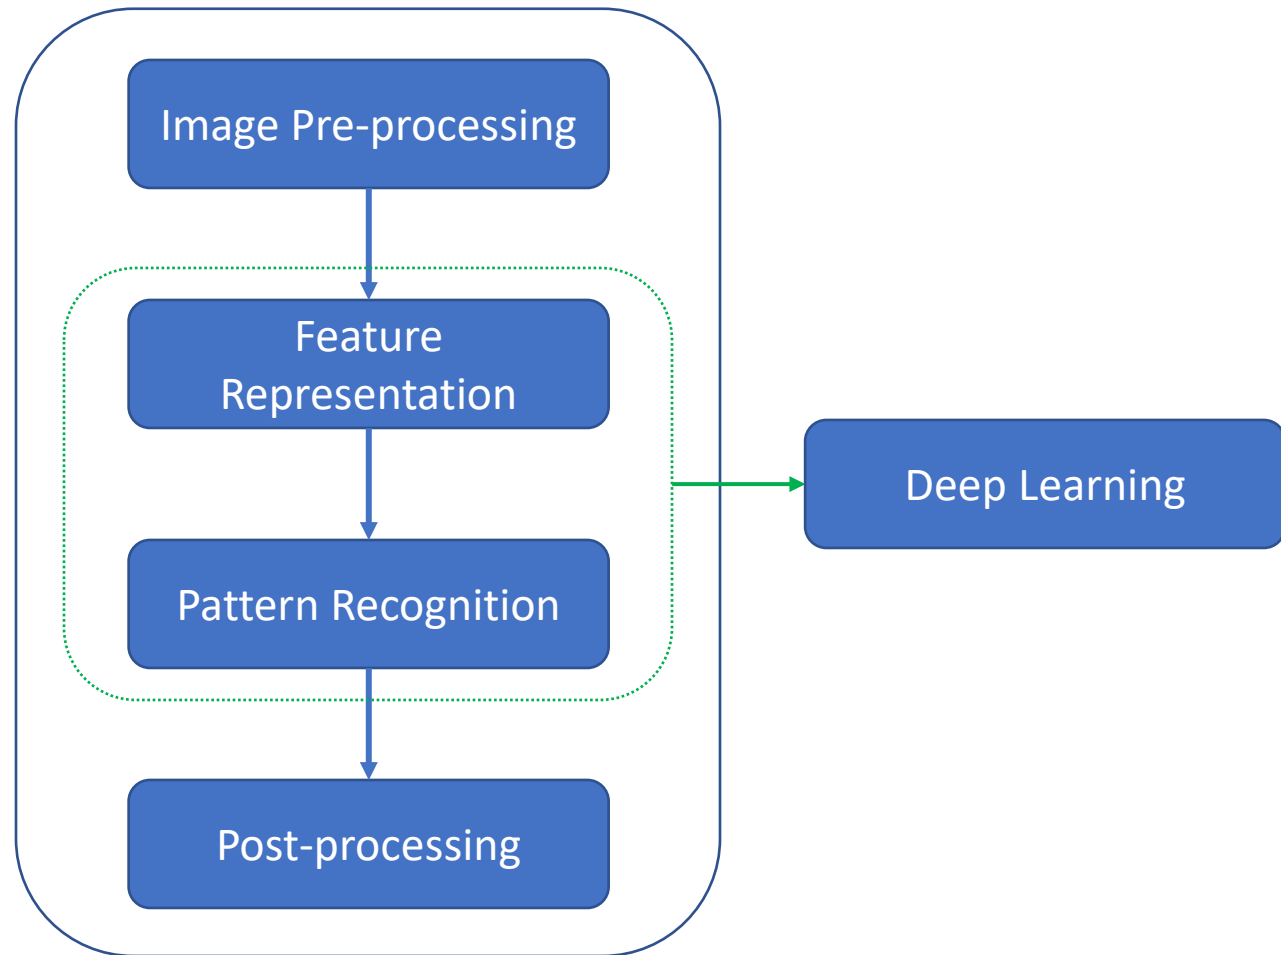
# Properties of Features

- Why not just use pixels values directly?
  - Pixel values change with light intensity, colour and direction
  - They also change with camera orientation
  - And they are highly redundant
- Repeatability (robustness)
  - Should be detectable at the same locations in different images despite changes in illumination and viewpoint
- Saliency (descriptiveness)
  - Similar salient points in different images should have similar features
- Compactness (efficiency)
  - Fewer features
  - Smaller features



# General Framework

Object detection  
Image segmentation  
Image classification  
Image retrieval  
Image stitching  
Object tracking  
...



# Feature Types

- Colour features
  - Colour histogram
  - Colour moments
- Texture features
  - Haralick texture features
  - Local binary patterns (LBP)
  - Scale-invariant feature transform (SIFT)
  - Texture feature encoding
- Shape features
  - Basic shape features
  - Shape context
  - Histogram of oriented gradients (HOG)

# Colour Features

- **Colour** is the simplest feature to compute, and is invariant to image scaling, translation and rotation.
- Example: colour-based image retrieval

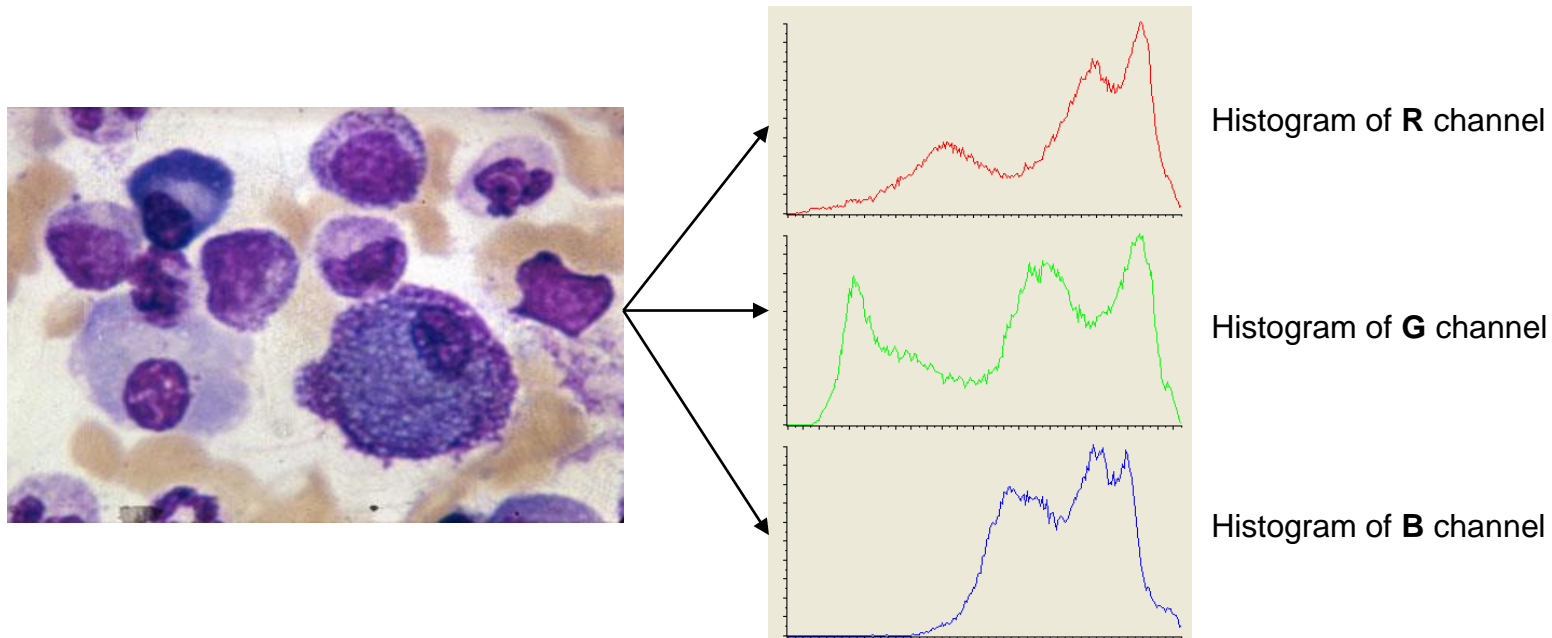


<http://labs.tineye.com/multicolr/>



# Colour Histogram

- Represent the global distribution of pixel colours in an image
  - Step 1: Construct a histogram for each colour channel (R, G, B)
  - Step 2: Concatenate the histograms (vectors) of all channels as the final feature vector



# Colour Moments

$f_{ij}$  is the value of the  $i$ -th colour component of pixel  $j$  and  $N$  is the number of pixels in the image

- Another way of representing colour distributions

- First-order moment

$$\mu_i = \frac{1}{N} \sum_{j=1}^N f_{ij}$$

(*mean*)

- Second-order moment

$$\sigma_i = \left( \frac{1}{N} \sum_{j=1}^N (f_{ij} - \mu_i)^2 \right)^{\frac{1}{2}}$$

(*variance*)

- Third-order moment

$$s_i = \left( \frac{1}{N} \sum_{j=1}^N (f_{ij} - \mu_i)^3 \right)^{\frac{1}{3}}$$

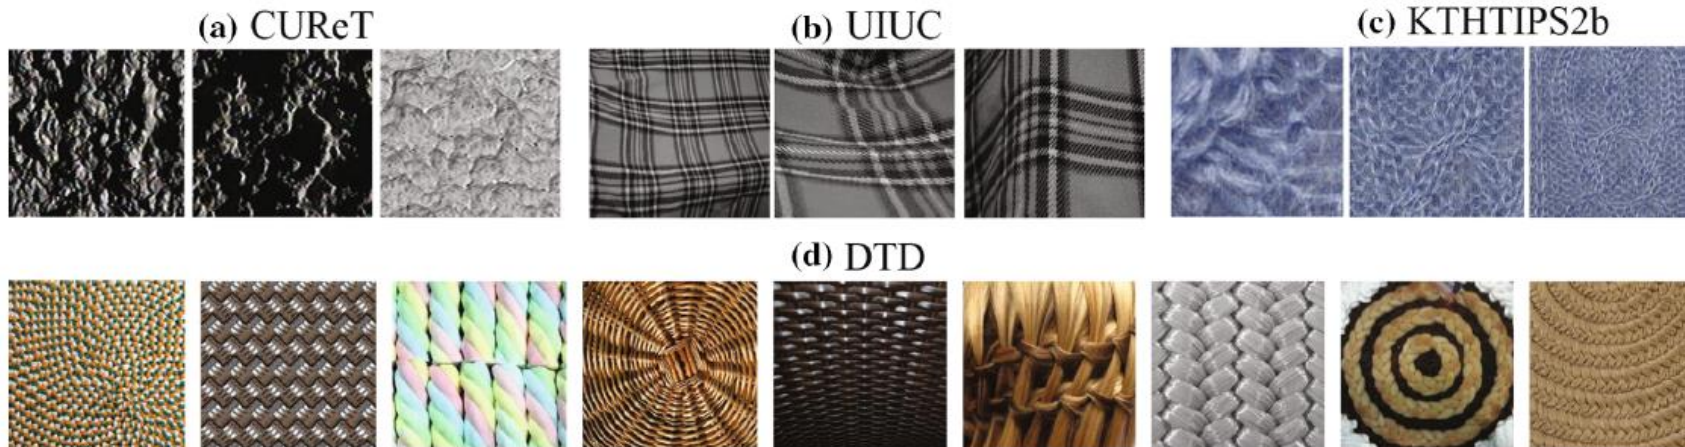
(*skewness*)

- Moments based representation of colour distributions

- Gives a feature vector of only 9 elements (for RGB images)
- Lower representation capability than the colour histogram

# Texture Features

- **Texture** is a powerful discriminating feature for identifying visual patterns with properties of homogeneity that cannot result from the presence of only a single color or intensity
- Example: texture classification



<https://arxiv.org/abs/1801.10324>



# Haralick Features

- Haralick features give an array of statistical descriptors of image patterns to capture the spatial relationship between neighbouring pixels, that is, textures
  - Step 1: Construct the gray-level co-occurrence matrix (GLCM)
  - Step 2: Compute the Haralick feature descriptors from the GLCM

610

IEEE TRANSACTIONS ON SYSTEMS, MAN, AND CYBERNETICS, VOL. SMC-3, NO. 6, NOVEMBER 1973

## Textural Features for Image Classification

ROBERT M. HARALICK, K. SHANMUGAM, AND ITS'HAK DINSTEIN

*Abstract*—Texture is one of the important characteristics used in identifying objects or regions of interest in an image, whether the image be a photomicrograph, an aerial photograph, or a satellite image. This paper describes some easily computable textural features based on gray-tone spatial dependancies, and illustrates their application in category-

array. If  $L_x = \{1, 2, \dots, N_x\}$  and  $L_y = \{1, 2, \dots, N_y\}$  are the  $X$  and  $Y$  spatial domains, then  $L_x \times L_y$  is the set of resolution cells and the digital image  $I$  is a function which assigns some gray-tone value  $G \in \{1, 2, \dots, N_g\}$  to each and every

# Haralick Features




Fig. 1 Resolution cells 1 and 5 are  $0^\circ$  (horizontal) nearest neighbors to resolution cell \*; resolution cells 2 and 6 are  $135^\circ$  nearest neighbors; resolution cells 3 and 7 are  $90^\circ$  nearest neighbors; and resolution cells 4 and 8 are  $45^\circ$  nearest neighbors to \*. (Note this information is purely spatial, and has nothing to do with gray-tone values.)

- Step 1: Construct the GLCMs
  - Given a distance  $d$  at an orientation angle  $\vartheta$ , then  $p_{(d, \vartheta)}(l_1, l_2)$ , being the  $(l_1, l_2)$  coefficient of the corresponding matrix  $\mathbf{P}_{(d, \vartheta)}$ , is the co-occurrence count or probability of going from a grey level  $l_1$  to another grey level  $l_2$  with an inter-sample spacing of  $d$  along the axis making an angle  $\vartheta$  with the x axis.
  - If the number of distinct gray levels in the quantized image is  $L$ , then the co-occurrence matrix  $\mathbf{P}$  will be of size  $L \times L$ .

0	0	1	1
0	0	1	1
0	2	2	2
2	2	3	3

original  $4 \times 4$  image

$$\mathbf{P}_{(1,0^\circ)} = \begin{bmatrix} 4 & 2 & 1 & 0 \\ 2 & 4 & 0 & 0 \\ 1 & 0 & 6 & 1 \\ 0 & 0 & 1 & 2 \end{bmatrix}$$

co-occurrence matrix construction

$$\mathbf{P}_{(1,135^\circ)} = \begin{bmatrix} 2 & 1 & 3 & 0 \\ 1 & 2 & 1 & 0 \\ 3 & 1 & 0 & 2 \\ 0 & 0 & 2 & 0 \end{bmatrix}$$

# Haralick Features

- Step 1: Construct the GLCMs
  - For computational efficiency, the number of gray levels ( $L$ ) can be reduced by binning (similar to histogram binning), e.g.  $L = 256/n$ , with  $n$  a constant factor.
  - Different co-occurrence matrices can be constructed by using various combinations of distance ( $d$ ) and angular directions ( $\vartheta$ ).
  - On their own, these co-occurrence matrices do not provide any measure of texture that can easily be used as texture descriptors.
  - The information in the co-occurrence matrices needs to be further extracted as a set of feature values => Haralick descriptors.



# Haralick Features

- Step 2: Compute the Haralick descriptors from the GLCMs
  - One set of Haralick descriptors for each GLCM corresponding to a particular distance ( $d$ ) and angular direction ( $\vartheta$ )

Angular Second Moment

$$\sum_i \sum_j p(i, j)^2$$

Contrast

$$\sum_{n=0}^{N_g-1} n^2 \left\{ \sum_{i=1}^{N_g} \sum_{j=1}^{N_g} p(i, j) \right\}, |i - j| = n$$

Correlation

$$\frac{\sum_i \sum_j (ij) p(i, j) - \mu_x \mu_y}{\sigma_x \sigma_y}$$

where  $\mu_x$ ,  $\mu_y$ ,  $\sigma_x$ , and  $\sigma_y$   
are the means and std. deviations  
of  $p_x$  and  $p_y$ , the partial probability  
density functions

Sum of Squares: Variance

$$\sum_i \sum_j (i - \mu)^2 p(i, j)$$

Inverse Difference Moment

$$\sum_i \sum_j \frac{1}{1 + (i - j)^2} p(i, j)$$

Sum Average

$$\sum_{i=2}^{2N_g} i p_{x+y}(i)$$

where  $x$  and  $y$  are the coordinates (row and  
column) of an entry in the co-occurrence matrix,  
and  $p_{x+y}(i)$  is the probability of co-occurrence  
matrix coordinates summing to  $x + y$

# Haralick Features

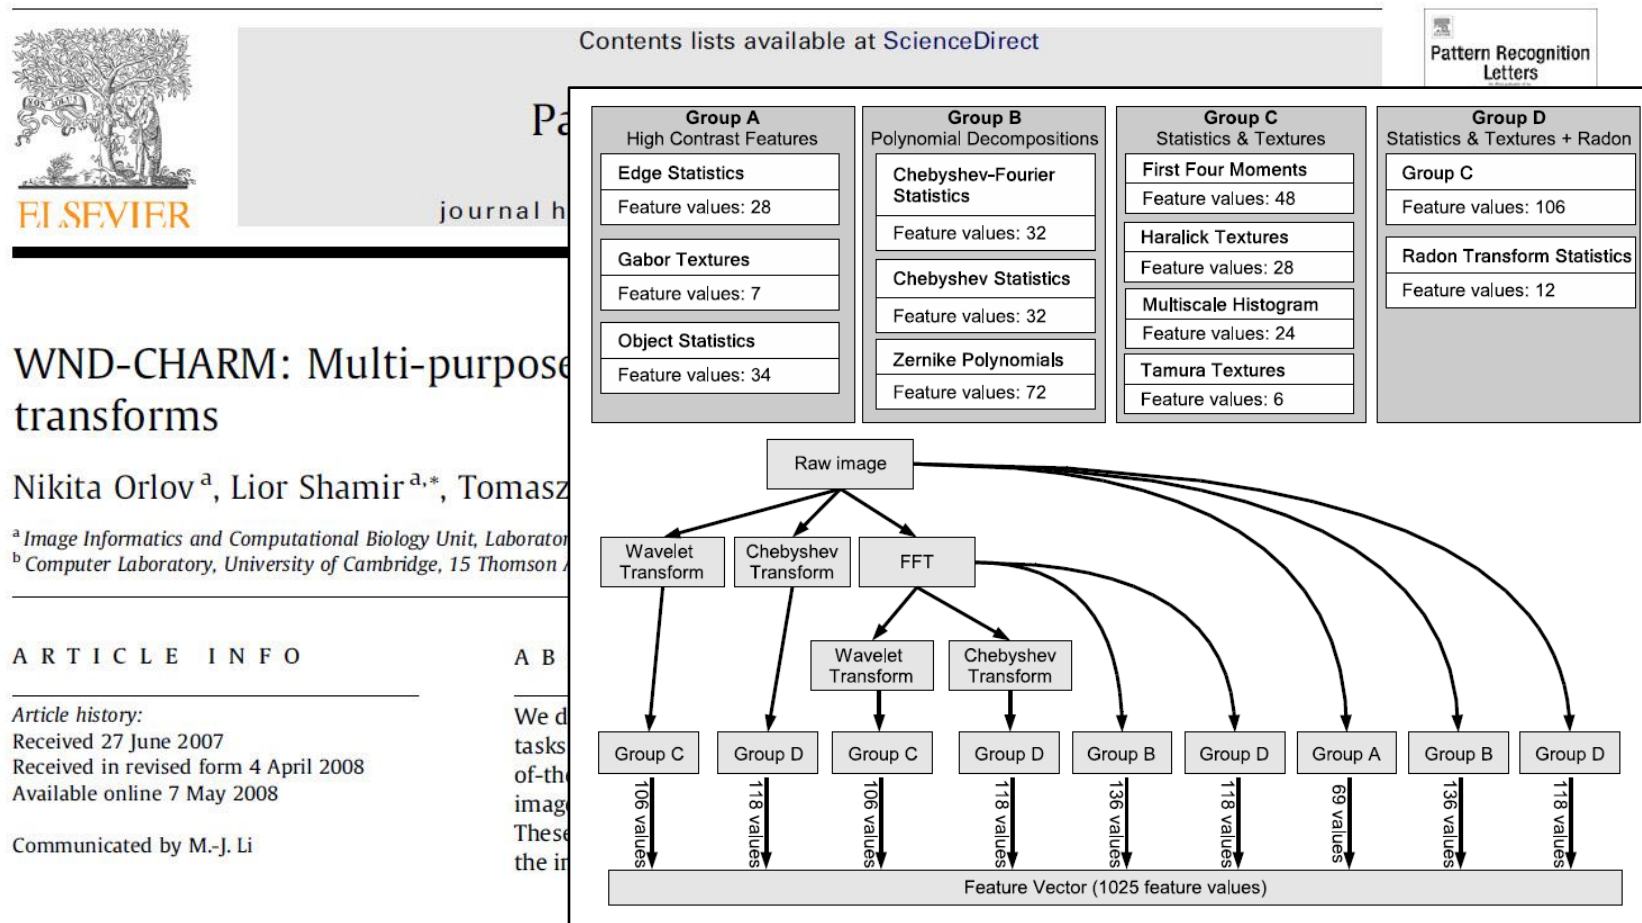
- Step 2: Compute the Haralick descriptors from the GLCMs
  - One set of Haralick descriptors for each GLCM corresponding to a particular distance ( $d$ ) and angular direction ( $\vartheta$ )

Sum Variance	$\sum_{i=2}^{2N_g} (i - f_8)^2 p_{x+y}(i)$
Sum Entropy	$-\sum_{i=2}^{2N_g} p_{x+y}(i) \log\{p_{x+y}(i)\} = f_8$
Entropy	$-\sum_i \sum_j p(i, j) \log(p(i, j))$
Difference Variance	$\sum_{i=0}^{N_g-1} i^2 p_{x-y}(i)$
Difference Entropy	$-\sum_{i=0}^{N_g-1} p_{x-y}(i) \log\{p_{x-y}(i)\}$
Info. Measure of Correlation 1	$\frac{HXY - HXY1}{\max\{HX, HY\}}$
Info. Measure of Correlation 2	$(1 - \exp[-2(HXY2 - HXY)])^{\frac{1}{2}}$ where $HXY = -\sum_i \sum_j p(i, j) \log(p(i, j))$ , $HX$ , $HY$ are the entropies of $p_x$ and $p_y$ , $HXY1 =$ $-\sum_i \sum_j p(i, j) \log\{p_x(i)p_y(j)\}$ $HXY2 =$ $-\sum_i \sum_j p_x(i)p_y(j) \log\{p_x(i)p_y(j)\}$
Max. Correlation Coeff.	Square root of the second largest eigenvalue of $\mathbf{Q}$ where $\mathbf{Q}(i, j) = \sum_k \frac{p(i, k)p(j, k)}{p_x(i)p_y(k)}$

# Application Example


Pattern Recognition Letters 29 (2008) 1684–1693

<https://doi.org/10.1016/j.patrec.2008.04.013>



# Application Example

- Commonly used nowadays in medical imaging studies due to its simplicity and interpretability



C. Jensen et al.

*Assessment of prostate cancer prognostic Gleason grade group using zonal-specific features extracted from biparametric MRI using a KNN classifier*

Journal of Applied Clinical Medical Physics, 2019


<https://doi.org/10.1002/acm2.12542>

1. Pre-processing
2. Extract Haralick, run-length, and histogram features from the region of interest
3. Feature selection
4. Classification using kNN



# Local Binary Patterns

- Describe the spatial structure of local image texture
  - Divide the image into cells of  $N \times N$  pixels (e.g.  $N = 16$  or  $32$ )
  - Compare each pixel in a cell to each of its 8 neighbouring pixels:  
If the centre pixel's value is greater than the neighbour's value, write "0", otherwise write "1"
  - This gives an 8-digit binary pattern per pixel after comparing with all 8 neighbouring pixels, representing a value in the range 0...255



# Local Binary Patterns

- Describe the spatial structure of local image texture (cont.)
  - Generate the histogram for all pixels in the cell, computing the frequency of each 8-digit binary number occurring in the cell
  - This gives a 256-bin histogram (the LBP feature vector)
  - Combine the histograms of all cells to obtain the image-level LBP feature descriptor


100	155	110	210
157	187	222	100
219	218	255	210
215	218	255	233



A histogram of 256 elements

# Local Binary Patterns

- LBP can be multi-resolution and rotation-invariant
  - Multi-resolution: varying the distance between the centre pixel and neighbouring pixels, and the number of neighbouring pixels



T. Ojala, M. Pietikainen, and T. Maenpaa, "Multiresolution gray-scale and rotation invariant texture classification with local binary patterns," IEEE Transactions on Pattern Analysis and Machine Intelligence 24(7):971-987, 2002. <https://doi.org/10.1109/TPAMI.2002.1017623>

# Local Binary Patterns

- LBP can be multi-resolution and rotation-invariant
  - Rotation-invariant: varying the way of constructing the 8-digit binary number, e.g. performing bitwise shift to derive the smallest number


Example:

1 1 1 1 0 0 0 0	= 240	}	15
1 1 1 0 0 0 0 1	= 225		
1 1 0 0 0 0 1 1	= 195		
1 0 0 0 0 1 1 1	= 135		
0 0 0 0 1 1 1 1	= 15		
0 0 0 1 1 1 1 0	= 30		
0 0 1 1 1 1 0 0	= 60		
0 1 1 1 1 0 0 0	= 120		

Note: not all patterns have 8 shifted variants (e.g. 11001100 has only 4)

# Local Binary Patterns

- LBP can be multi-resolution and rotation-invariant
  - Rotation-invariant: varying the way of constructing the 8-digit binary number, e.g. performing bitwise shift to derive the smallest number  
=> this reduces the LBP feature dimension from 256 to 36




T. Ojala, M. Pietikainen, and T. Maenpaa, "Multiresolution gray-scale and rotation invariant texture classification with local binary patterns," IEEE Transactions on Pattern Analysis and Machine Intelligence 24(7):971-987, 2002. <https://doi.org/10.1109/TPAMI.2002.1017623>



# Application Example

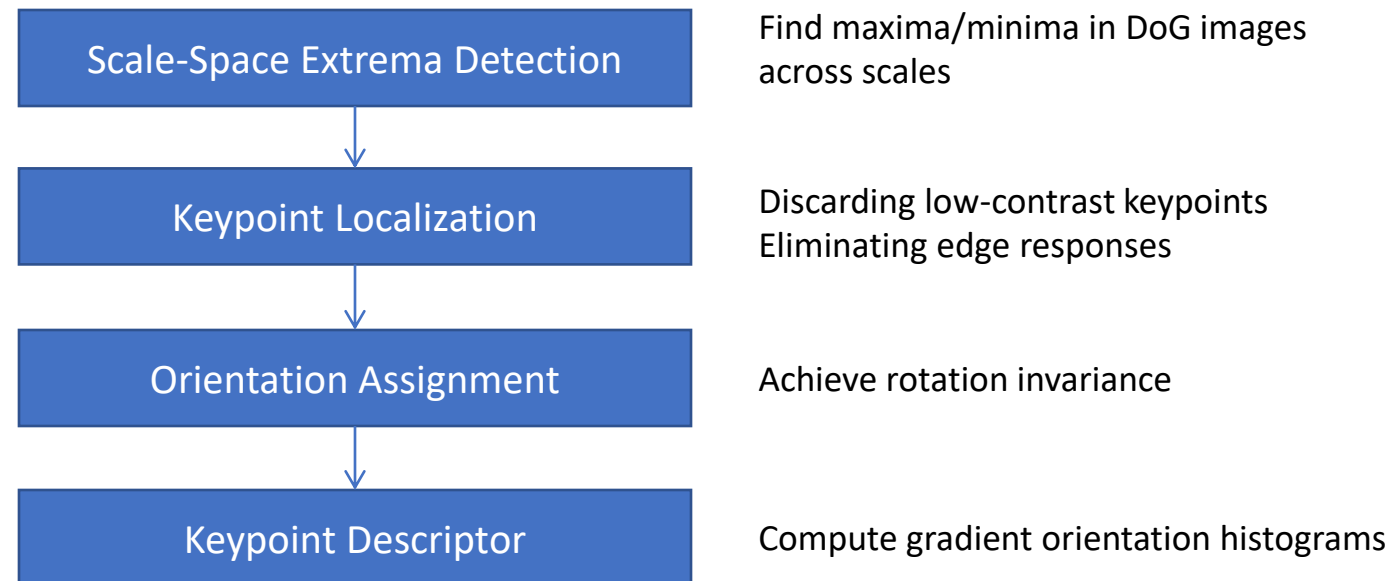
- Texture classification

$P,R$	$LBP_{P,R}$	
	BINS	RESULT
8,1	10	88.2
16,2	18	98.5
24,3	26	99.1
8,1+16,2	10+18	99.0
8,1+24,3	10+26	99.6
16,2+24,3	18+26	99.0
8,1+16,2+24,3	10+18+26	99.1




# Scale-Invariant Feature Transform

- SIFT feature describes the texture features in a localised region around a keypoint
- SIFT descriptor is invariant to uniform scaling, orientation, and partially invariant to affine distortion and illumination changes



# SIFT Extrema Detection


- Detect maxima and minima in the scale space of the image



D. G. Lowe, "Distinctive image features from scale-invariant keypoints," Int. J. Comput. Vis. 60(2):91-110, November 2004. <https://doi.org/10.1023/B:VISI.0000029664.99615.94>

# SIFT Keypoint Localization


- Improve and reduce the set of found keypoints
  - Use 3D quadratic fitting in scale-space to get subpixel optima
  - Reject low-contrast and edge points using Hessian analysis



Initial keypoints from  
scale-space optima





Keypoints after rejecting  
low-contrast points



Final keypoints after  
rejecting edge points

# SIFT Orientation Assignment


- Estimate keypoint orientation using local gradient vectors
  - Make an orientation histogram of local gradient vectors
  - Find the dominant orientations from the peaks of the histogram





# SIFT Keypoint Descriptor

- 4 x 4 array of gradient histogram weighted by magnitude
- 8 bins in gradient orientation histogram
- Total 8 x 4 x 4 array = 128 dimensions




# Application Example

- Image matching



# Application Example


- Image matching
  - Find SIFT keypoints





# Application Example

- Image matching
  - Find best matching between SIFT keypoints




# Descriptor Matching

- Nearest Neighbour Distance Ratio

$$NNDR = \frac{d_1}{d_2} = \frac{\|D_A - D_B\|}{\|D_A - D_C\|}$$


- $d_1$  is the distance to the first nearest neighbour
- $d_2$  is the distance to the second nearest neighbour
- Neighbours in feature space






# Application Example

- Image stitching



# Application Example

- Image stitching
  - Find SIFT keypoints and feature correspondences



# Application Example

- Image stitching
  - Find the right spatial transformation



# Transformations



original



translation



rotation



scale



shear



perspective

# Transformations

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

Scale

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & \alpha_x \\ \alpha_y & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

Shear

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos \Theta & -\sin \Theta \\ \sin \Theta & \cos \Theta \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

Rotate

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Affine

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

Projective



# Fitting and Alignment

- Least-squares (LS) fitting of corresponding keypoints  $(\mathbf{x}_i, \mathbf{x}'_i)$

$$E_{LS} = \sum_i \|\mathbf{r}_i\|^2 = \sum_i \|f(\mathbf{x}_i; \mathbf{p}) - \mathbf{x}'_i\|^2$$

where  $\mathbf{p}$  are the parameters of the transformation  $f$


$$\begin{aligned} \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} &= \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \quad \Rightarrow \quad \begin{bmatrix} x & y & 0 & 0 & 1 & 0 \\ 0 & 0 & x & y & 0 & 1 \\ & & \dots & & & \\ & & \dots & & & \end{bmatrix} \begin{bmatrix} a \\ b \\ d \\ e \\ c \\ f \end{bmatrix} = \begin{bmatrix} x' \\ y' \\ \vdots \end{bmatrix} \\ & \quad \Downarrow \\ \mathbf{p} &= [\mathbf{A}^T \mathbf{A}]^{-1} \mathbf{A}^T \mathbf{b} \quad \Leftarrow \quad \mathbf{A} \mathbf{p} = \mathbf{b} \end{aligned}$$

# Fitting and Alignment

- RANdom SAmple Consensus (RANSAC) fitting
  - Least-squares fitting is hampered by outliers
  - Some kind of outlier detection and rejection is needed
  - Better use a subset of the data and check inlier agreement
  - RANSAC does this in a iterative way to find the optimum

# Fitting and Alignment

- RANSAC



(line fitting example)


Algorithm:

1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence

# Fitting and Alignment

- RANSAC



(line fitting example)


Algorithm:

1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence

# Fitting and Alignment

- RANSAC




Algorithm:

1. **Sample** (randomly) the number of points required to fit the model
  2. **Solve** for model parameters using samples
  3. **Score** by the fraction of inliers within a preset threshold of the model
- Repeat** 1-3 until the best model is found with high confidence

# Fitting and Alignment

- RANSAC



Algorithm:


1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence



# Fitting and Alignment


- RANSAC



Algorithm:

1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model



**Repeat** 1-3 until the best model is found with high confidence



# Fitting and Alignment

- Given matched points A and B, estimate the translation


$$\begin{bmatrix} x_i^B \\ y_i^B \end{bmatrix} = \begin{bmatrix} x_i^A \\ y_i^A \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$



# Alignment by RANSAC

1. Sample a set of matching points (1 pair)
2. Solve for transformation parameters
3. Score parameters with number of inliers
4. Repeat steps 1-3 N times

$$\begin{bmatrix} x_i^B \\ y_i^B \end{bmatrix} = \begin{bmatrix} x_i^A \\ y_i^A \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$



# Application Example

- SIFT-based texture classification – how to do this?

bread



cracker

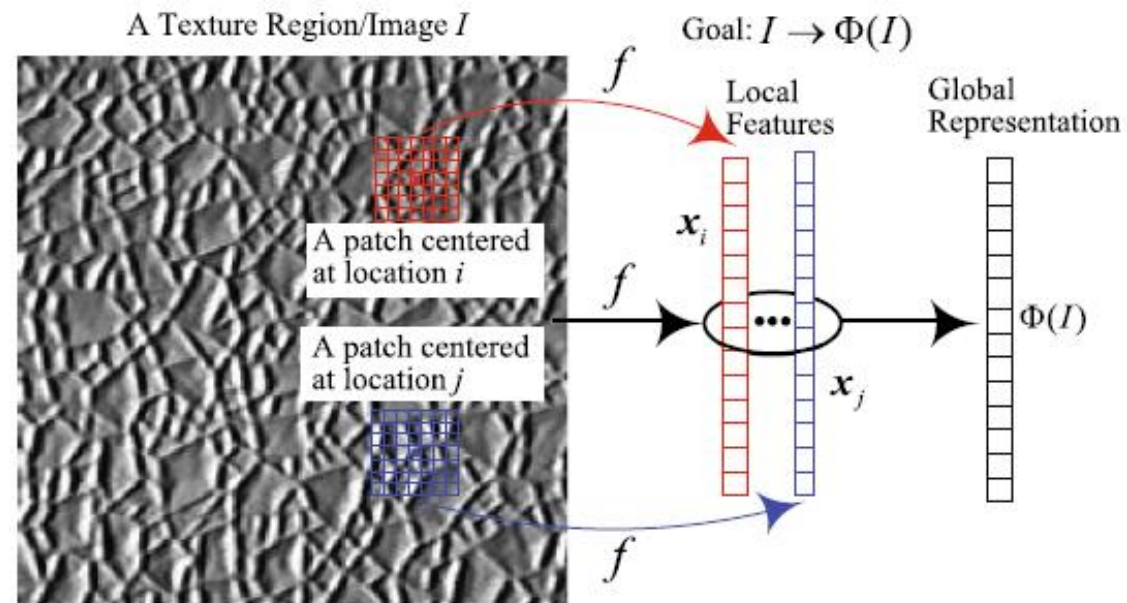


**Problem:** the number of SIFT keypoints (and thus the number of SIFT feature descriptors) may vary highly between images



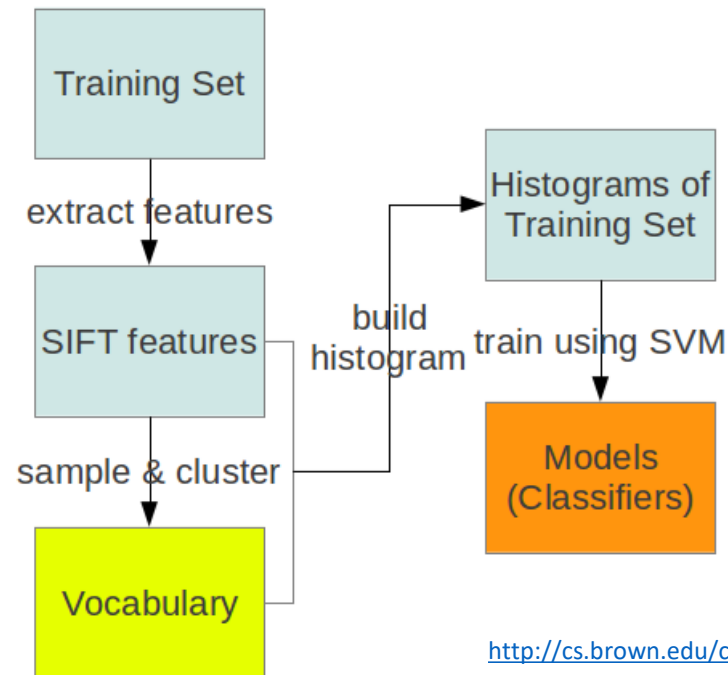
# Feature Encoding

- Global encoding of local SIFT features
  - Integrate the local features (SIFT keypoint descriptors) of an image into a global vector to represent the whole image



# Feature Encoding

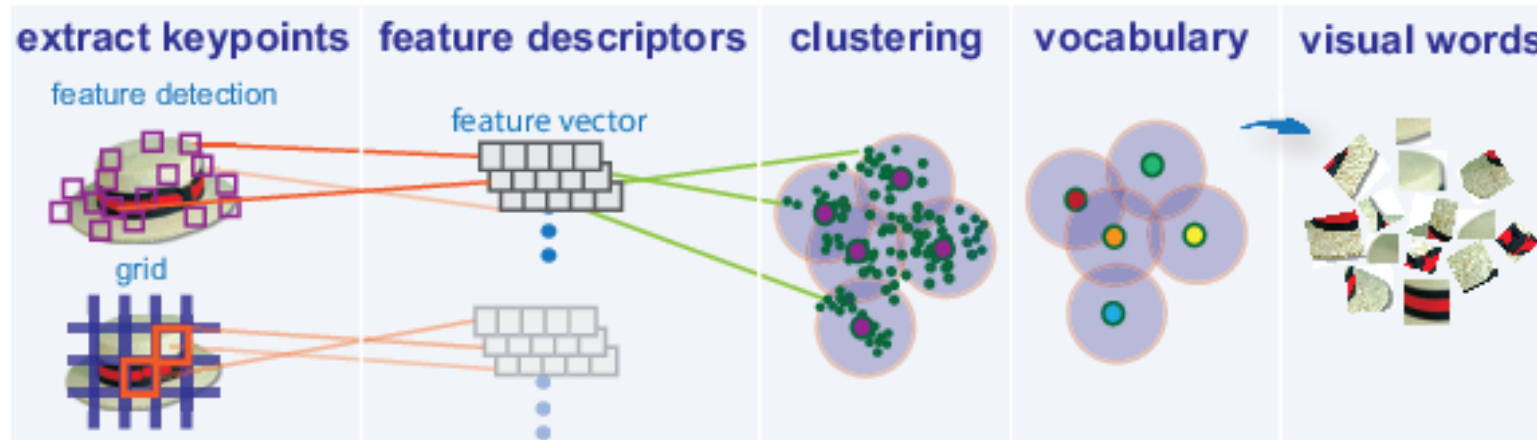
- Most popular method: Bag-of-Words (BoW)
  - The variable number of local image features are encoded into a fixed-dimensional histogram to represent each image



<http://cs.brown.edu/courses/cs143/2011/results/proj3/hangsu/>


# Feature Encoding

- Bag-of-Words (BoW) – step 1
  - Create the vocabulary from the set of local descriptors (SIFT keypoint descriptors) extracted from the training data
  - This vocabulary represents the categories of local descriptors




# Feature Encoding

- Bag-of-Words (BoW) – step 1
  - Main technique used to create the vocabulary: *k-means clustering*
  - k-means clustering is one of the simplest and most popular unsupervised learning approaches that perform automatic clustering (partitioning) of the training data into multiple categories




# Feature Encoding

- Bag-of-Words (BoW) – step 1
  - K-means clustering:
    - Initialize:  $k$  cluster centres, typically randomly
    - Iterate: 1) Assign data (feature vectors) to the closest cluster (Euclidean distance)  
2) Update cluster centres as the mean of the data samples in each cluster
    - Terminate: When converged or the number of iterations reaches the maximum





# K-Means Clustering




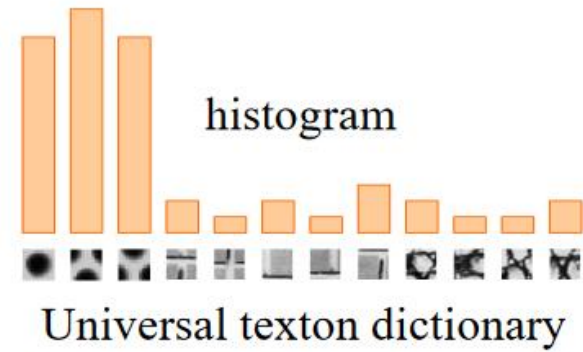




# Feature Encoding

- Bag-of-Words (BoW) – step 2
  - The cluster centres are the “visual words” which form the “vocabulary” that is used to represent an image
  - An individual local feature descriptor (e.g. SIFT keypoint descriptor) is assigned to one visual word with the smallest distance



# Feature Encoding



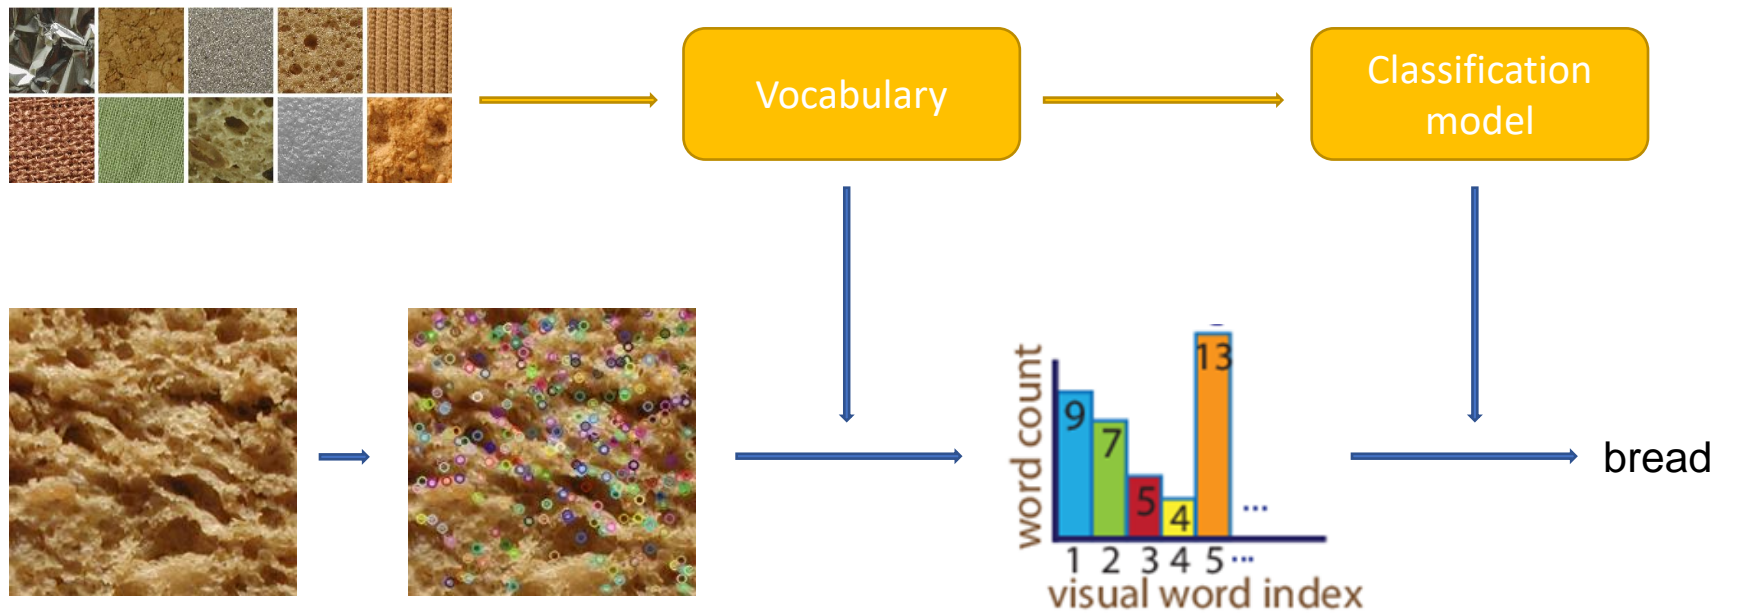
# Feature Encoding

- Bag-of-Words (BoW) – step 2
  - For an image the number of local feature descriptors assigned to each visual word is computed
  - The numbers are concatenated into a vector which forms the BoW representation of the image



# Application Example

- SIFT-based texture classification



1. SIFT feature extraction

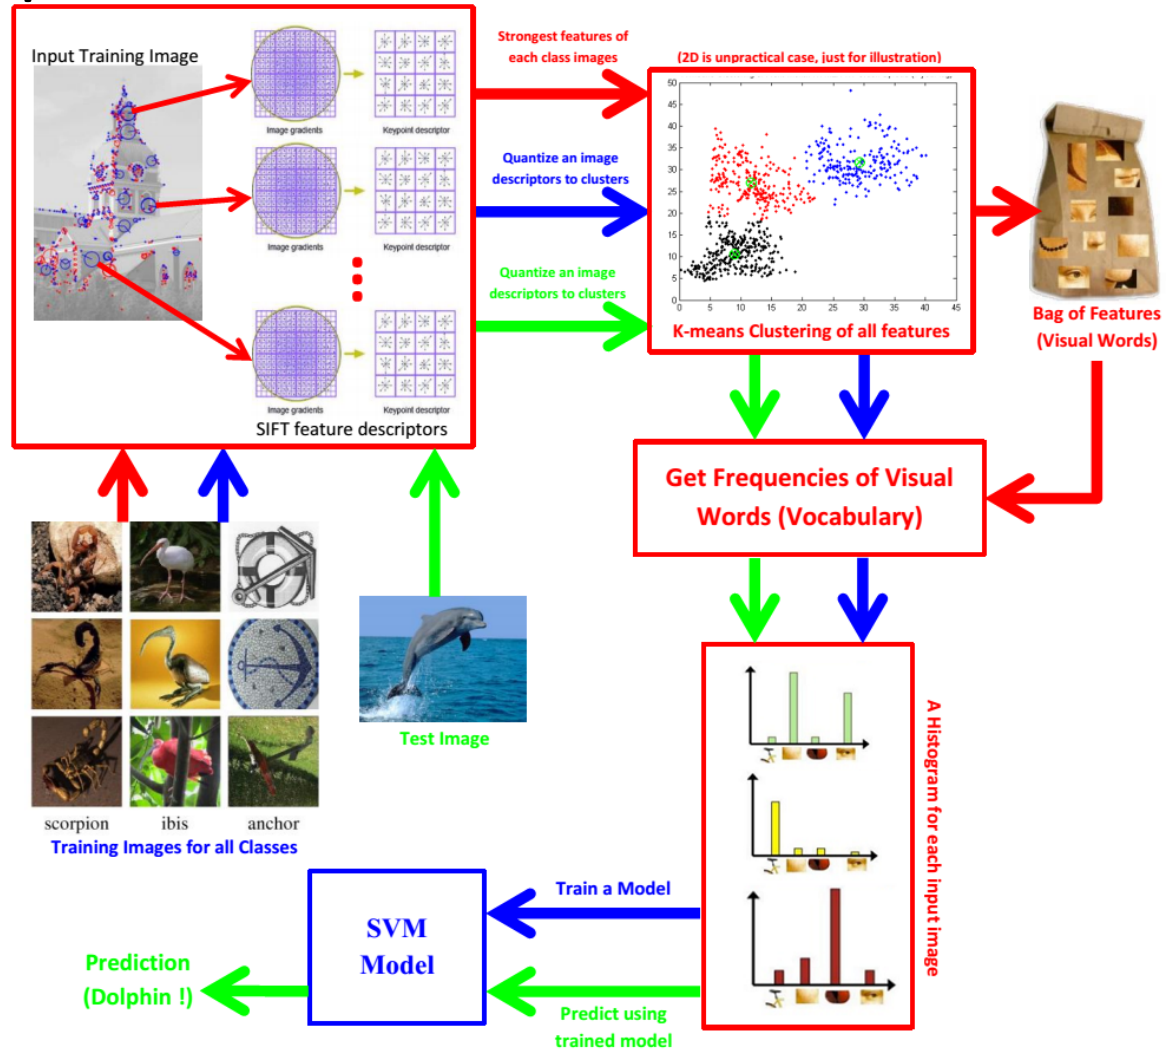
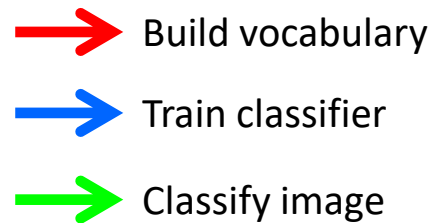
2. BoW encoding

3. Classification



# Application Example

- SIFT-based texture classification



<http://heraqi.blogspot.com/2017/03/BoW.html>

# Feature Encoding

- Local features can be other types of features, not just SIFT
  - LBP, SURF, BRIEF, ORB
- There are also more advanced techniques than BoW
  - VLAD, Fisher Vector
- A very good source of additional information is VLFeat.org
  - <http://www.vlfeat.org/>

# Summary

- Feature representation is essential in solving almost all types of computer vision problems
- Most commonly used image features:
  - Colour features (Part 1)
    - Colour moments and histogram
  - Texture features (Part 1)
    - Haralick, LBP, SIFT
  - Shape features (Part 2)
    - Basic, shape context, HOG

# Summary

- Other techniques described (Part 1)
  - Descriptor matching
  - Feature encoding (Bag-of-Words)
  - k-means clustering
  - Alignment and RANSAC
  - Spatial transformations
- To be discussed (Part 2)
  - Shape features
  - Shape matching
  - Sliding window detection

# References and Acknowledgements

- Szeliski, Chapter 4 (in particular Sections 4.1.1 to 4.1.3 and 4.3.2), Chapter 6 (in particular Sections 6.1.1 to 6.1.4)
- Some content are extracted from the above resource, James Hays slides, slides from Michael A. Wirth and Cordelia Schmit
- L. Liu et al., [From BoW to CNN: two decades of texture representation for texture classification](#), International Journal of Computer Vision, 2019
- And other resources as indicated by the hyperlinks