

COMP9517: Computer Vision

Image Segmentation

Part 2

Outline




- Basic segmentation methods
 - Thresholding
 - K-means clustering
 - Feature extraction and classification
- More sophisticated segmentation methods
 - Region splitting and merging
 - Watershed segmentation
 - Maximally stable extremal regions
 - Mean-shift algorithm
 - Supapixel segmentation
 - Conditional random field
 - **Active contour segmentation**
 - **Level-set segmentation**

Outline

- **Region Representation**
 - Labelled images (the most commonly used)
 - Overlays
 - Boundary coding
 - Quad trees
 - Property tables
- **Evaluating segmentation methods**
 - Quantitative evaluation metrics
 - Receiver operating characteristic



Active Contour Segmentation

- A contour based approach to object segmentation
 - Aims to locate object boundaries in images by curve fitting
 - Represents the curve by a set of control points and interpolation
 - Iteratively moves the control points to fit the curve to the object
 - Uses **image**, **smoothness** and **user-guidance forces** along curve



Active Contours

- Examples of implementations:
 - Snakes
 - Level sets




- Active contours can also be used in a wide variety of object-tracking applications
- Based on global optimisation

Global Optimisation

- formulate the goals of the desired transformation (e.g. segmentation) using some optimisation criterion, and then infer the solution that best meets this criterion
- Problem: finding a smooth surface that passes through a set of measured data points
 - ill-posed: many possible surfaces can fit this data
 - ill-conditioned: small changes in the input can sometimes lead to large changes in the fit
 - inverse problems: recover the unknown function form which data point were sampled
- therefore, regularisation is needed
 - to fit models to data that badly underconstrained the solution space

Global Optimisation

- in order to quantify what it means to find a smooth solution, we can define a norm on the solution space
- the **derivative** is a measure of how a function changes as its input changes.



- a derivative can be thought of as how much one quantity is changing in response to changes in some other quantity
- for example, the derivative of the position of a moving object with respect to time is the object's instantaneous velocity

$$m = \frac{\Delta f(x)}{\Delta x} = \frac{f(x+h) - f(x)}{(x+h) - (x)} = \frac{f(x+h) - f(x)}{h}.$$

Global Optimisation

- For one dimensional function $f(x)$, to find a smooth solution, we can define a **norm** on the solution space:
 - integrate the squared first derivative of the function

$$\varepsilon_1 = \int f_x^2(x) dx$$

- integrate the squared second derivative

$$\varepsilon_2 = \int f_{xx}^2(x) dx$$

- For two dimensions(e.g. for images), the corresponding smoothness functions are (partial derivative)

$$\varepsilon_1 = \int f_x^2(x, y) + f_y^2(x, y) dx dy = \int \|\nabla f(x, y)\|^2 dx dy$$


$$\varepsilon_2 = \int f_{xx}^2(x, y) + 2f_{xy}^2(x, y) + f_{yy}^2(x, y) dx dy$$

Global Optimization

- We also require a data term (or data penalty)
- For scattered data interpolation, the data term measures distance between function $f(x,y)$ and a set of data points $d_i = d(x_i, y_i)$
- By adding the norm and the data term, we get a global energy term that can be minimized

Active Contour Segmentation

- A well-known implementation is called the **snakes algorithm**
 - Smoothly follows high intensity gradients at the object boundary
 - Bridges areas of noise or missing gradients using smooth interpolation



Snakes

- Snakes are a two-dimensional generalisation of the 1-D energy minimising splines

$$\mathcal{E}_{\text{int}} = \int \alpha(s) \|f_s(s)\|^2 + \beta(s) \|f_{ss}(s)\|^2 ds$$

- discretised version of the energy function

$$E_{\text{int}} = \int \alpha(s) \|f(i+1) - f(i)\|^2 / h^2 + \beta(s) \|f(i+1) - 2f(i) + f(i-1)\|^2 / h^4$$

- additional external image-based and constraint-based potentials

$$\mathcal{E}_{\text{image}} = \omega_{\text{line}} \mathcal{E}_{\text{line}} + \omega_{\text{edge}} \mathcal{E}_{\text{edge}} + \omega_{\text{term}} \mathcal{E}_{\text{term}}$$

Snakes

- In practice, only the edge term is used
 - directly proportional to the image gradients

$$E_{edge} = \sum_i -\|\nabla I(f(i))\|^2$$

- a smoothed version of the image Laplacian

$$E_{edge} = \sum_i -\|(G_\delta \circ \nabla^2 I)(f(i))\|^2$$

- a distance map to the edges

- user-placed constraints


$$E_{spring} = k_i \|f(i) - d(i)\|^2$$

Snakes

- Put them together

$$E = E_{\text{int}} + E_{\text{image}} + E_{\text{spring}}$$


- Let E to be zero, the energy E can be iteratively minimised by moving the curve



(a)



(b)




Active Contour Segmentation

- Active contours / snakes are parametric models
 - Explicit representations of the object boundaries
 - Typically requires manual interaction to initialize the curve
 - It is challenging to change the topology of the curve as it evolves
 - Curve reparameterization may be required for big shape changes
- Level-set methods have become more popular alternatives
 - Implicit representations of the object boundaries
 - Boundaries defined by the zero-set of a higher dimensional function
 - Level-set function evolves to make the zero-set fit and track objects
 - Easily accommodates topological changes in object shape
 - Computationally more demanding than active contours

Level Set Segmentation


- Contour evolution over time (iterative)



https://en.wikipedia.org/wiki/Level-set_method

Level Set Segmentation

- Examples of level-set based segmentation



Level Set Segmentation

- A well-known implementation is the Chan-Vese model


Minimize the energy functional

$$E(C) = \mu \cdot \text{length}(C) + \nu \cdot \text{area}(C)$$

$$+ \lambda_1 \int_{\text{Inside}(C)} |u_0 - c_1|^2 dx dy$$

$$+ \lambda_2 \int_{\text{Outside}(C)} |u_0 - c_2|^2 dx dy$$

<https://doi.org/10.1109/83.902291>





Region Representation

- The outputs from segmentation methods need to be represented and stored
- Methods to represent regions:
 - Labelled images (the most commonly used)
 - Overlays
 - Boundary coding
 - Quad trees
 - Property tables


Region Representation

- Labelled images
 - assign each detected region a unique identifier (an integer)



Region Representation

- Labelled images
 - create an image where all the pixels of a region will have its unique identifier as their pixel value



0	0	0	2	2
0	0	1	1	2
3	1	1	1	2
3	3	1	3	2
3	3	3	3	3

Region Representation

- Overlays
 - overlaying some colour or colours on top of the original image




Region Representation


- Boundary Coding

- represents regions by their boundaries in a data structure instead of an image


- a linear list of the border pixels
 - chain code
 - polygon approximation



original curve




chain code links



encoding scheme

100076543532

chain code representation



polygonal approximation

Region Representation

- Boundary Coding
 - encoding schema
 - start point
 - chain code links
 - chain code representation
 - polygon approximation

Region Representation


- Boundary Coding
 - encoding schema

– start point


– character blocks

– character representation

– polygon approximation



8-N



4-N

Region Representation

- Boundary Coding


- encoding schema

- start Point

- chain code links

- chain code representation



- polygon approximation



Region Representation


- Boundary Coding

- encoding schema
- start Point
- chain code links
- chain code representation
- polygon approximation




Region Representation

- Boundary Coding
 - encoding schema
 - start Point
 - chain code links
 - chain code representation
 - polygon approximation




Chain Code:

2223446756701





Region Representation

- Quadtrees
 - a **quadtree** is a tree data structure in which each internal node has exactly four children
 - quadtrees are most often used to partition a two dimensional space by recursively subdividing it into four quadrants or regions




Region Representation

- Quadtrees
 - space-saving representation encoding the whole region
 - each region of interest would be represented by a quadtree structure.




Region Representation

- Quadtrees
 - encoding Schema
 - split the image
 - create the tree





Region Representation

- Quadtrees
 - encoding Schema
 - split the image
 - create the tree



Region Representation



- Quadtrees
 - encoding Schema
 - split the image
 - create the tree



Region Representation


- Quadtrees

- each node of a quadtree representing a square region in the image has one of three labels: *Full*, *Empty* or *Mixed*
- only mixed nodes have child nodes



Region Representation

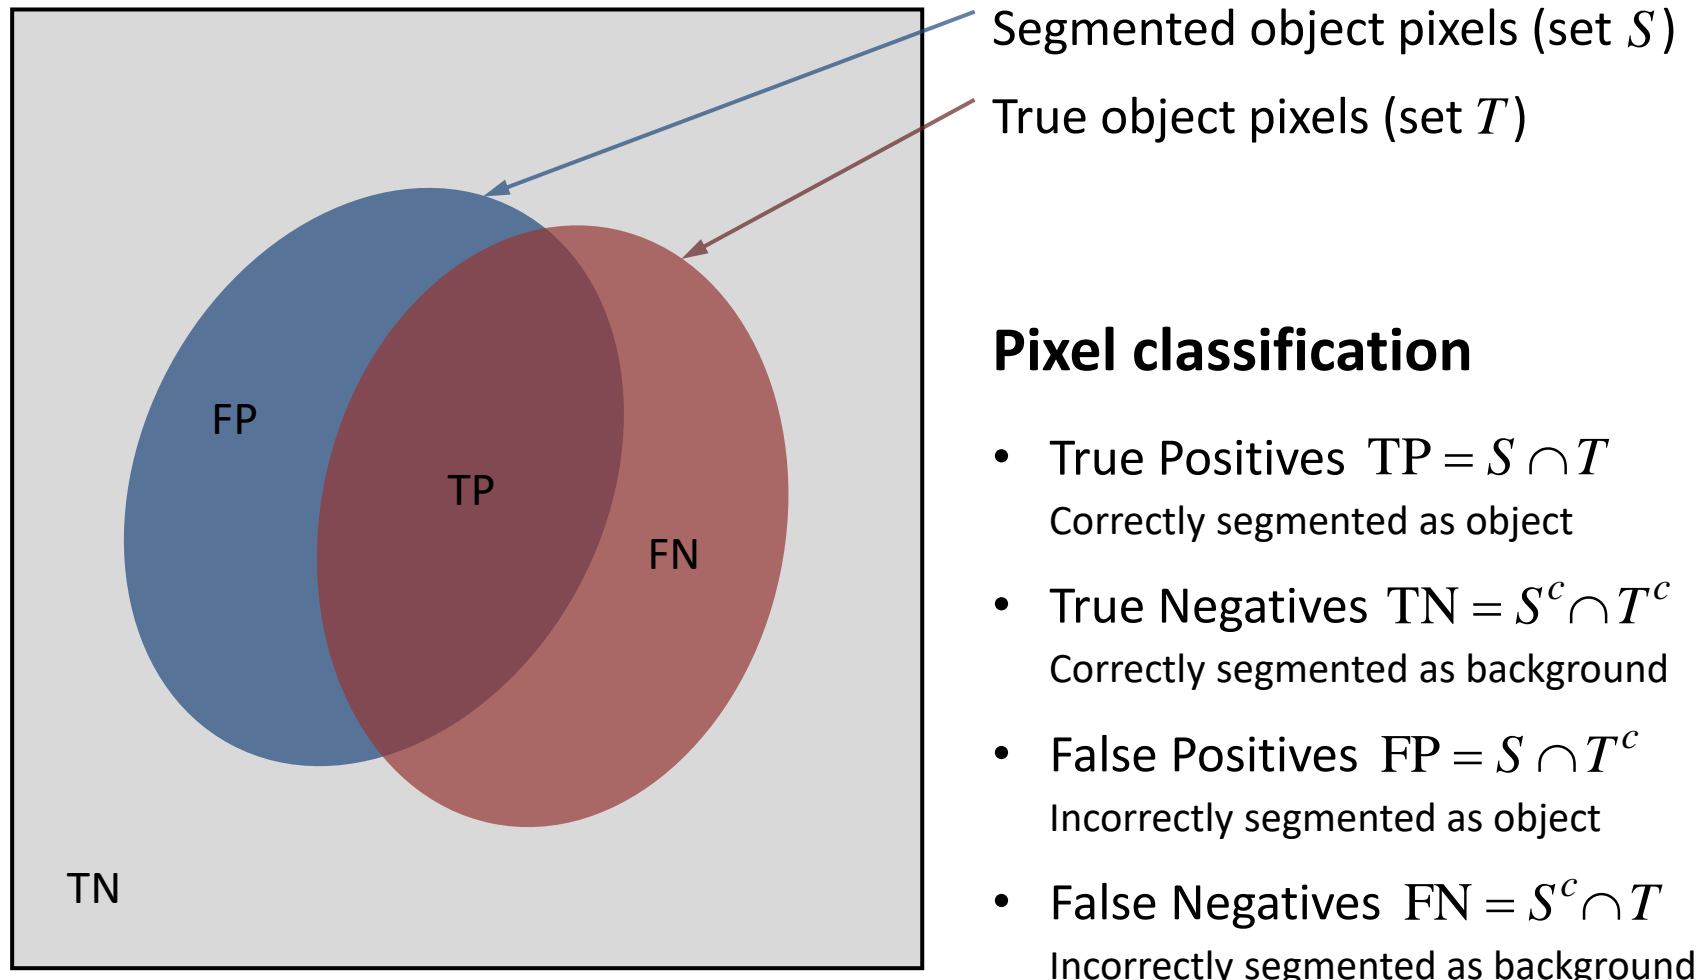
- Property Tables
 - represent a region by its extracted properties rather than by its pixels
 - properties can be the size, shape, intensity, colour or texture of the region




Region	Size	colour	...
0	5	Red	
1	6	Blue	
2	5	Green	
3	9	Grey	

0	0	0	2	2
0	0	1	1	2
3	1	1	1	2
3	3	1	3	2
3	3	3	3	3

Segmentation Evaluation Metrics




Evaluation Metrics



- Sensitivity (= true-positive rate)

Fraction of the true object that is correctly segmented

$$\text{TPR} = \frac{|\text{TP}|}{|\text{TP}| + |\text{FN}|}$$



- Specificity (= true-negative rate)

Fraction of the true background that is correctly segmented


$$\text{TNR} = \frac{|\text{TN}|}{|\text{TN}| + |\text{FP}|}$$

Evaluation Metrics

- Receiver operating characteristic (ROC) of a method

Plot of the true-positive rate (sensitivity) versus the false-positive rate (one minus the specificity) of a method as a function of its free parameter(s)

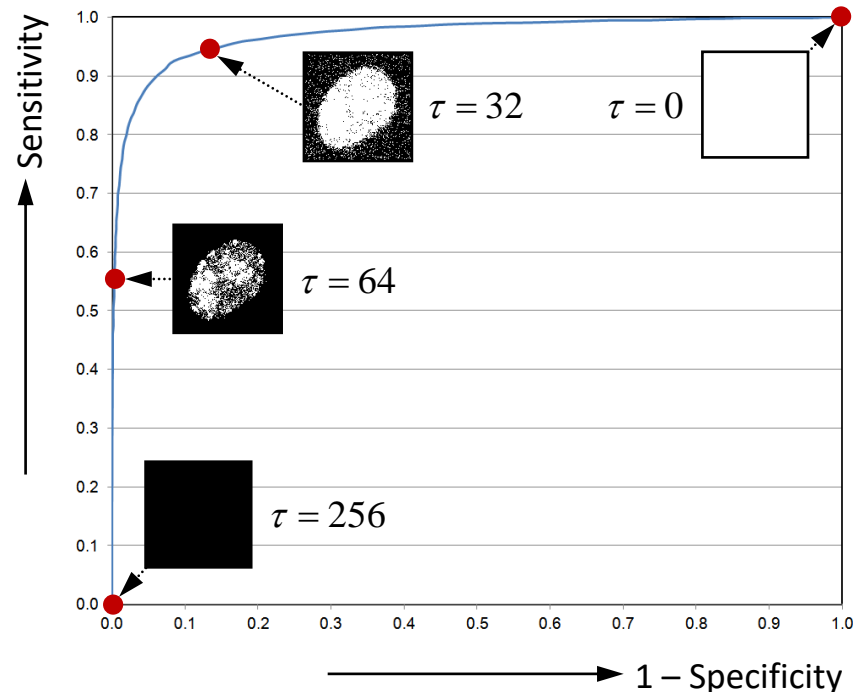
Example: Thresholding



Image


Truth

Compute the sensitivity versus the specificity for all possible intensity thresholds τ and plot the results



Evaluation Metrics

- Comparing the performance of methods by ROC analysis



Area under the curve (AUC)

Worst = 0.0

0.1

0.2

0.3

0.4

Random = 0.5

0.6

0.7


0.8

0.9

Best = 1.0

Higher AUC = better method


Evaluation Metrics



- Precision (= positive predictive value)

Fraction of the segmented object that is correctly segmented

$$P = \frac{|TP|}{|TP| + |FP|}$$



- Recall (= sensitivity)

Fraction of the true object that is correctly segmented


$$R = \frac{|TP|}{|TP| + |FN|}$$

- F-measure

Harmonic mean of precision and recall


$$F = 2 \cdot \frac{R \cdot P}{R + P}$$

Evaluation Metrics



- Jaccard similarity coefficient (JSC)
Fraction of the union of the segmented object and the true object that is correctly segmented

$$\text{JSC} = \frac{|S \cap T|}{|S \cup T|} = \frac{|\text{TP}|}{|\text{FP}| + |\text{TP}| + |\text{FN}|}$$



- Dice similarity coefficient (DSC)
Fraction of the segmented object set joined with the true object set that is correctly segmented

$$\text{DSC} = 2 \frac{|S \cap T|}{|S| + |T|} = \frac{2|\text{TP}|}{|\text{FP}| + 2|\text{TP}| + |\text{FN}|}$$

References and Acknowledgements

- Chapter 3 & 5 of Szeliski 2010
- Shapiro and Stockman 2001
- Some images drawn from Szeliski 2010
- Some images drawn from papers as indicated